Another, Doubtless Very Different, Bill James

Alex Groce
Northern Arizona University
Flagstaff, USA
agroce@gmail.com

Abstract

Software engineering metrics and baseball metrics have a lot in common; we oughta do something about that.

CCS Concepts: • **Software and its engineering** → **Dynamic analysis**; **Software testing and debugging**.

Keywords: software metrics

ACM Reference Format:
Alex Groce. 2026. Another, Doubtless Very Different, Bill James. In *Proceedings of the 2024 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software (Onward! '26), October 3–9, 2026, Oakland, CA, USA.* ACM, New York, NY, USA, 3 pages. https://doi.org/10.1145/3689492.3689810

1 Introduction

Harlan D. Mills (1919-1996) is the namesake of the Harlan D. Mills Award given by the IEEE Computer Society to recognize "researchers and practitioners who have demonstrated long-standing, sustained, and meaningful contributions to the theory and practice of the information sciences, focusing on contributions to the practice of software engineering through the application of sound theory." First given in 1999, recipients include such luminaries as David Parnas (the first winner), Elaine Weyuker, Bertrand Meyer, the Cousots, Gerard Holzmann, Pamela Zave, Gail Murphy, Mark Harman, Matt Dwyer, David Harel, and, most recently, Andreas Zeller. Mills' idea of the Chief Programmer Team influenced Brooks' proposals for how to avoid software disaster in *The Mythical Man-Month*; his Cleanroom conception informs every approach that is focused on defect-prevention (from static analysis to aggressive unit testing to mutation testing) rather than removing defects found in the field; Mills was also a major force in introducing both statistics (viewing tests as statistical experiments, which arguably also is a foundational insight for all of random testing and fuzzing) and formal specification (thus a contributor to the lines of thought that have given the field Alloy, TLA+, Dafny, and other treasures). In short, he was a heckuva big shot in our field, though (other than having one of the top awards in the field named after him) I think he's one of the least talked-about major visionaries (one name for the Mills Award, confusingly implied as being started in 1996, and perhaps by ICSE/the ACM, in Wikipedia, is the "Harlan Mills Practical Visionary Prize") we have.

I know the rough outlines of the above details because I am a software engineering history buff. I know that Mills was a command pilot for the US Army Air Corps in WWII, and rewrote the B-24 pilot training program because he thought it was garbage, and that he worked at the Institute for Advanced Study alongside Einstein, Godel, and von Neumann, because Harlan Mills and his brother Eldon are the most prominent characters in Chapter 4 of Alan Schwarz's book, *The Numbers Game: Baseball's Lifelong Fascination with Statistics* [**?**]. Because, you see, besides a passion for applied mathematics and operational research and software engineering, Harlan Mills had a passion for baseball. Harlan and Eldon Mills, in 1968, founded a company called Computer Research in Sports. Their goal was to go beyond the (much-beloved: one could argue that batting average, which we'll talk about below, drove the country crazy a few times in the years since baseball became America's pastime) statistics baseball fans already knew to statistics that truly measured how a player contributed to winning a game. They were among many high-powered intellects (including multiple Nobel Prize winners and the popular Harvard evolutionary-biology maverick, public intellectual, and superb writer Stephen Jay Gould) pioneering an effort to truly understand baseball that is documented in Schwarz's engaging, in-depth, look at baseball's love of statistics and how those statistics have developed from the earliest, alien-looking, box-scores.

The connection between baseball statistics and software engineering, however, goes beyond the curiosity that Mills was a pioneer in both software engineering and baseball stats-savvy. The obvious connection is that the stats revolution in baseball was in substantial part tied to the computer revolution; nothing like https://www.retrosheet.org/ could plausibly exist in a non-computerized world, and certainly looking for trends or even just finding individual facts in such a database (play-by-play accounts of over 100,000 pro

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
Onward! '26, October 3–9, 2026, Oakland, CA, USA
© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 979-8-4007-1215-9/24/10
https://doi.org/10.1145/3689492.3689810

baseball games) would be impossible even if a warehouse of filing cabinets had the relevant data. Schwarz's book shows how computers were essential to the growth of the modern baseball stats world, and takes us back to when baseball work was done by "stealing" off-time on hulking company and government mainframes. The deeper connection, however, essay addresses, is that baseball was, for a very long time, tied to extremely traditional, and beloved statistics, but even stodgy, change-averse professional baseball front and back offices eventually realized the limitations of these statistics and moved on to more sophisticated, nuanced, statistics that truly reflected how to win the game. That story: famous, widely-used numbers, lacking in proven operational relevance, should sound familliar to students of software engineering.

2 Build a Better Mousetrap/Metric

To develop the theme of the similarities (data is easy to gather, hard to understand; a lot of people who like futzing with spreadsheets and coding up formulas are involved; there's a lot of money and a lot of (to stay polite, this is a family paper!) "BS" involved) and differences (baseball is an extraordinarily structured thing, and the binary outcome of almost every event (did someone make an out? did someone score? did the team win?) is obvious, while in software engineering figuring out if a computer program works and was well-made turns out to be extremely annoyingly hard unless you wait twenty years and see if anyone is using it and how much money it made) between baseball and software engineering would require considerable argument and evidence. But perhaps we can quickly see what I mean by the move from traditional stats to ones that mean more by looking at a simple example from software testing (we could also compare LOC produced to modern software productivity measures that are less simpleminded, like DevOps' Change Failure Rate [3]).

One way to determine if a program is well-tested is to take the test suite for the program and look at the statement coverage: the percent of all statements in the whole program that are executed at least once by the test suite. Obviously if a test suite fails to execute most of the code in the program it is ostensibly testing, it is likely a pretty bad test suite! Similarly, if a test suite runs 90% of the code in the program, it's presumably at least doing something non-trivial right.

In baseball, for a very very long time, the king of statistics was batting average. What is batting average? Even if you know little about baseball, even if you're a European or from Mars [1], you may well know that in baseball, a batter is thrown balls by the pitcher and tries to hit those balls, and then (when the ball is hit) run to a series of three bases, before returning to the place from which the ball is initially hit (home base). Batting average looks at how often a player hits the ball to advance to first, second, third, or home base. Batting average is not nonsense! The most sophisticated baseball statisticians tend to agree players celebrated for their batting averages, such as Ty Cobb, Rogers Hornsby, and Ted Williams, were superb, top-of-the-line players; the top players by batting average now, like Aaron Judge, tend to be household names for a good reason. Batting average, like code coverage, is a *meaningful number*.

However, there are some problems with statement coverage, and code coverage in general. For one thing, maybe you only run a really important line of code once in the whole test suite, and that in a rather boring and uninteresting situation where the things the line of code is supposed to do are not seriously stressed (it's a complex multiplication, but one factor is zero, for example). Or, worse yet, you may run lots of code lots of times, but not check anything much about whether the program did what it was supposed to do other than "not crash." There are lots of ways a program can go wrong other than outright crashing. Mutation score [2, 4] asks "what do we really care about in testing? We want to find bugs. Statement coverage is an awful roundabout way to measure finding bugs, since bugs aren't even mentioned here. Let's see how a test suite does at finding bugs." In mutation testing, a large number of essentially random bugs are actually injected into a program, and the score is the percent of the fake bugs the test suite was able to find. It's not perfect, but if you think about it, it completely removes the two obvious problems with statement coverage mentioned above.

Similarly, the way batting average is defined, it doesn't look at every time a player comes up to home plate to try to hit the ball. The numerator is "times the player hit the ball and got to a base" but the denominator is "times the player came up to bat, minus times the pitcher did such a bad job the player was allowed to take a base without hitting the ball, and the batter took proper advantage of this by not swinging at bad pitches (or times the pitcher feared the player so much this was intentionally done to keep the bat off the ball), times the pitcher hit the player with the ball, times the catcher did something weird, times the player intentionally hit the ball in a way that tends not to let you reach base but has another game-winning-relevant purpose, etc. etc. etc.). On-base-percentage, a revolution in the 1980s in understanding how to value hitters, counted most everything, and looked (essentially) at "of times this player came up to bat, how often did this player end up at least standing on a base." Like mutation score, it focused on the *more* essential (not getting out before reaching a base, finding bugs) rather than the easily and obviously measured (hitting the ball, running the code). The typical non-baseball-fan probably still knows batting average, not on-base-percentage; the typical software developer (who *should* be more like a baseball fan than someone who isn't sure if Babe Ruth and Willy Mays

[1] The best kid's book on baseball is Walter R. Brooks' (Brooks is also the guy who invented Mr. Ed) *Freddy and the Baseball Team from Mars*, by the way).

played at the same time or not), alas, probably is still more familliar with code coverage than mutation score, but the game is changing, and highly visible adoption of mutation at, for example, Google [5] and Meta [1] suggest a revolution in this metric is underway.

3 Why It Isn't Easy

The above suggests that improving software engineering metrics is easy: just do what the baseball people did, and look at whether the most widely used numbers are measuring the most important things.

References

[1] Moritz Beller, Chu-Pan Wong, Johannes Bader, Andrew Scott, Mateusz Machalica, Satish Chandra, and Erik Meijer. 2021. What It Would Take to Use Mutation Testing in Industry - A Study at Facebook. In *International Conference on Software Engineering: Software Engineering in Practice (ICSE '21)*. IEEE, 268–277. https://doi.org/10.1109/ICSE-SEIP52600.2021.00036

[2] R. A. DeMillo, R. J. Lipton, and F. G. Sayward. 1978. Hints on Test Data Selection: Help for the Practicing Programmer. *IEEE Computer* 11, 4 (Apr 1978), 34–41. https://doi.org/10.1109/C-M.1978.218136

[3] DevOps Research and Assessment (DORA). 2014–2024. State of DevOps Reports (various years). https://dora.dev/. Defines "change failure rate" as the percentage of production changes that result in degraded service or require remediation (e.g., rollback, hotfix, patch).

[4] Alex Groce, Mohammad Amin Alipour, and Rahul Gopinath. 2014. Coverage and Its Discontents. In *Proceedings of the 2014 ACM International Symposium on New Ideas, New Paradigms, and Reflections on Programming & Software* (Portland, Oregon, USA) *(Onward 2014)*. Association for Computing Machinery, New York, NY, USA, 255–268. https://doi.org/10.1145/2661136.2661157

[5] Goran Petrovic and Marko Ivankovic. 2018. State of Mutation Testing at Google. In *International Conference on Software Engineering: Software Engineering in Practice (ICSE '18)*. 163–171. https://doi.org/10.1145/3183519.3183521